Chi2: Feature Selection and Discretization of Numeric Attributes

Huan Liu and Rudy Setiono
Department of Information Systems and Computer Science
National University of Singapore
Kent Ridge, Singapore 0511
{liuh,rudys}@iscs.nus.sg

Abstract

Discretization can turn numeric attributes into discrete ones. Feature selection can eliminate some irrelevant attributes. This paper describes Chi2, a simple and general algorithm that uses the χ^2 statistic to discretize numeric attributes repeatedly until some inconsistencies are found in the data, and achieves feature selection via discretization. The empirical results demonstrate that Chi2 is effective in feature selection and discretization of numeric and ordinal attributes.

1 Introduction

Feature selection is a task to select the minimum number of attributes needed to represent the data accurately. By using relevant features, classification algorithms can in general improve their predictive accuracy, shorten the learning period, and result in the simpler concepts. There are abundant feature selection algorithms [5]. Our work adopts an approach that selects a subset of the original attributes since it not only has the above virtues, but also serves as an indicator on what kind of data (along those selected features) should be collected. The feature selection algorithms can be further divided based on the data types they operate on. The basic two types of data are nominal (e.g., attribute *color* may have values of red, green, yellow) and ordinal (e.g., attribute *winning position* can have values of 1, 2, and 3, or attribute *salary* can have 22345.00, 46543.89, etc. as its values). Many feature selection algorithms [1, 3, 5] are shown to work effectively on discrete data or even more strictly, on binary data (and/or binary class value). In order to deal with numeric attributes, a common practice for those algorithms is to discretize the data before conducting feature selection. This paper provides a way to select features directly from numeric attributes while discretizing them. Numeric data are very common in real world problems. However, many classification algorithms require that the training data contain only discrete attributes, and some would work better on discretized or binarized data [2, 4]. If those numeric data can be *automatically* transformed into discrete ones, these classification algorithms would be readily at our disposal. Chi2 is our effort towards this goal: discretize the numeric attributes as well as select features among them.

The problem this work tackles is as follows: there are data sets with numeric attributes, some are irrelevant and the range of each numeric attribute could be very wide; find an algorithm that can automatically discretize the numeric attributes as well as remove those irrelevant ones.

This work stems from Kerber's ChiMerge [4] which is designed to discretize numeric attributes based on the χ^2 statistic. ChiMerge consists of an initialization step and a bottom-up merging process, where intervals are continuously merged until a termination condition, which is determined by a significance level α (set manually), is met. It is an improvement from the most obvious simple methods such as *equal-width-intervals* or *equal-frequency-intervals*. Instead of defining a width or frequency threshold (which is not easy until scrutinizing each attribute and knowing what it is), ChiMerge requires α to be specified. Nevertheless, too big or too small an α will over- or under-discretize an attribute. An extreme example of under-discretization is the continuous attribute itself. Over-discretization will introduce many inconsistencies[1] nonexistent before, thus change the characteristics of the data. In short, it is not easy to find a proper α for ChiMerge. It is thereby ideal to let the data determine what value α should take. This leads to Phase 1 of Chi2. Naturally, if the discretization continues without generating more inconsistencies than in the original data, it is possible that some attributes will be discretized into one interval only. Hence, they can be removed.

2 Chi2 Algorithm

The Chi2 algorithm (summarized below) is based on the χ^2 statistic, and consists of two phases. In the first phase, it begins with a high significance level (sigLevel), e.g., 0.5, for all numeric attributes for discretization. Each attribute is sorted according to its values. Then the following is performed: 1. calculate the χ^2 value as in equation (1) for every pair of adjacent intervals (at the beginning, each pattern is put into its own interval that contains only one value of an attribute); 2. merge the pair of adjacent intervals with the lowest χ^2 value. Merging continues until all pairs of intervals have χ^2 values exceeding the parameter determined by sigLevel (initially, 0.5, its

[1] By inconsistency we mean that two patterns are the same, but classified into different categories.

1082-3409/95 $04.00 © 1995 IEEE

corresponding χ^2 value is 0.455 if the degree of freedom is 1, more below). The above process is repeated with a decreased sigLevel until an inconsistency rate, δ is exceeded in the discretized data. Phase 1 is, as a matter of fact, a generalized version of ChiMerge of Kerber [4]. Instead of specifying a χ^2 threshold, Chi2 wraps up ChiMerge with a loop that automatically increments the χ^2 threshold (decrementing sigLevel). A consistency checking is also introduced as a stopping criterion in order to guarantee that the discretized data set accurately represents the original one. With these two new features, Chi2 automatically determines a proper χ^2 threshold that keeps the fidelity of the original data.

Phase 2 is a finer process of Phase 1. Starting with sigLevel0 determined in Phase 1, each attribute i is associated with a sigLevel[i], and takes turns for merging. Consistency checking is conducted after each attribute's merging. If the inconsistency rate is not exceeded, sigLevel[i] is decremented for attribute i's next round of merging; otherwise attribute i will not be involved in further merging. This process is continued until no attribute's values can be merged. At the end of Phase 2, if an attribute is merged to only one value, it simply means that this attribute is not relevant in representing the original data set. As a result, when discretization ends, feature selection is accomplished.

Chi2 Algorithm:

```
Phase 1:
  set sigLevel = .5;
  do while (InConsistency(data) < δ) {
    for each numeric attribute {
      Sort(attribute, data);
      chi-sq-initialization(attribute, data);
      do {
        chi-sq-calculation(attribute,data)
      } while (Merge(data))
    }
    sigLevel0 = sigLevel;
    sigLevel = decreSigLevel(sigLevel);
  }
Phase 2:
  set all sigLvl[i] = sigLevel0 for attribute i;
  do until no-attribute-can-be-merged {
    for each attribute i that can be merged {
      Sort(attribute, data);
      chi-sq-initialization(attribute, data);
      do {
        chi-sq-calculation(attribute,data)
      } while (Merge(data))
      if (InConsistency(data) < δ)
        sigLvl[i] = decreSigLevel(sigLvl[i]);
      else
        attribute i cannot be merged;
    }
  }
```

The formula for computing the χ^2 value is:

$$\chi^2 = \sum_{i=1}^{2}\sum_{j=1}^{k}\frac{(A_{ij}-E_{ij})^2}{E_{ij}} \quad (1)$$

where:
k = number of (no.) classes,
A_{ij} = no. patterns in the ith interval, jth class,
R_i = no. patterns in the ith interval = $\sum_{j=1}^{k} A_{ij}$,
C_j = no. patterns in the jth class = $\sum_{i=1}^{2} A_{ij}$,
N = total no. patterns = $\sum_{i=1}^{2} R_i$,
E_{ij} = expected frequency of $A_{ij} = R_i * C_j/N$.
If either R_i or C_j is 0, E_{ij} is set to 0.1. The degree of freedom of the χ^2 statistic is one less the number of classes.

3 Experiments

Two sets of experiments are conducted. In the first set of experiments, we want to establish that 1. Chi2 helps improve predictive accuracy; and 2. Chi2 properly and effectively discretizes data as well as eliminates some irrelevant attributes. C4.5 [8] (an extension of ID3 [7]) is used to verify the effectiveness of Chi2. The reasons for our choice are 1. C4.5 (or ID3) works well for many problems and is well known, thus requiring no further description; and 2. C4.5 selects relevant features by itself in tree branching so it can be used as a benchmark, as in [5, 9, 1], to verify the effects of Chi2. In the second set of experiments, we have a closer examination of Chi2's ability of discretization and feature selection by introducing a synthetic data set and adding noise attributes to the existing data set. Through these more controlled data sets, we can better understand how effective Chi2 is.

3.1 Real data

Three data sets used in experiments are Iris, Wisconsin Breast Cancer and Heart Disease[2]. They have different types of attributes. The Iris data are of continuous attributes, the breast cancer data are of ordinal discrete ones, and the heart disease data have mixed attributes (numeric and discrete).

3.2 Controlled data

Two extra data sets are designed to test if noise attributes can be removed. One is synthetic, the other is the Iris data added with noise attributes.

The synthetic data consists of 600 items and is described by four attributes among which only one attribute determines each item's class label. The values, v_1 of attribute A_1 are generated from a uniform distribution between the lower bound (L = 0) and the upper bound (U = 75), each item's class label is determined as follows: $v_0 < 25 \rightarrow$ class 1, $v_0 < 50 \rightarrow$ class 2, $v_0 < 75 \rightarrow$ class 3. Then we add noise attributes A_2, A_3, and A_4. The values of A_2 are generated from a normal distribution with $\mu = U/2$ (i.e. 37.5) and $\sigma = \mu/3$. The values of A_3 are generated

[2]They are all obtained from the University of California-Irvine machine learning repository via anonymous ftp to *ics.uci.edu*.

Int	Class Freq			χ^2	Int	Class Freq			χ^2
4.4	3	0	0	0.20	6.0	0	2	1	1.43
4.6	2	0	0	0.20	6.1	0	0	1	0.54
4.7	1	0	0	0.20	6.2	0	1	2	0.14
4.8	3	0	0	1.97	6.3	0	2	3	0.14
4.9	1	0	1	2.62	6.4	0	1	2	0.16
5.0	3	1	0	0.10	6.5	0	1	3	1.97
5.1	3	1	0	0.70	6.6	0	1	0	2.50
5.2	2	0	0	0.20	6.7	0	1	4	0.73
5.3	1	0	0	0.41	6.8	0	1	1	0.10
5.4	3	1	0	1.32	6.9	0	1	1	0.85
5.5	1	2	0	1.66	7.0	0	1	0	2.10
5.6	0	4	0	2.50	7.1	0	0	1	0.20
5.7	1	1	0	1.28	7.4	0	0	1	0.20
5.8	1	2	2	1.20	7.7	0	0	2	
5.9	0	1	0	0.54					

Table 1: The initial intervals, class frequencies, and χ^2 values for sepal-length.

from two normal distributions with $\mu = U/3$ (i.e. 25), $\mu = 2 * U/3$ (i.e. 50) and $\sigma = \mu/3$ respectively, 300 items each distribution. The values of A_4 are generated from a uniform distribution.

The second data is a modified version of Iris data. Four noise attributes A_5, A_6, A_7 and A_8 are added to the Iris training data corresponding to the four original attributes. The values of each noise attribute are determined by a normal distribution with $\mu = ave$ and $\sigma = (max - min)/6$, where ave is the average value of, max and min are the maximum and minimum values of the original attribute. The choice of σ is to approximate $\mu/3$ if the corresponding original attribute is of uniform distribution. Now there are total eight attributes. The number of patterns used is 75.

3.3 Example

In this section, some steps of Chi2 processing for the Iris data are shown to demonstrate the behavior of Chi2. Table 1 shows the intervals, class frequencies, and χ^2 values of sepal-length after the initialization in Phase 1. The results for sepal-length after Phase 1 and Phase 2 are shown in Table 2. An inconsistency rate $\delta = 5\%$ is allowed in the experiment, that means up to 3 (75x0.05) inconsistencies are acceptable. Phase 1 stops at sigLevel = 0.2, $\chi^2 = 3.22$. That means the next sigLevel (0.1) will introduce more inconsistencies. When Phase 2 terminates, the values of both sepal-length and sepal-width are merged into one value, so they can be removed; and attributes petal-length and petal-width are discretized into four discrete values each. With the χ^2 threshold 3.22, for example, six discrete values are needed for attribute sepal-length: $< 4.4 \rightarrow 0$, $< 4.9 \rightarrow 1$, ..., $< 6.1 \rightarrow 4$, and $\geq 6.1 \rightarrow 5$. The last one reads if a numeric value is greater than and equal to 6.1, it is quantized to 5.

3.4 Empirical results on real data

First we show that after discretization, the number of attributes decreases for the three data sets (in Figure 1). For the Iris data, the number of attributes is

Int	Class Freq			χ^2
4.4	9	0	0	5.05
4.9	1	0	1	8.11
5.0	12	3	0	13.64
5.5	3	12	3	14.23
6.1	0	10	21	

(a)

Int	Class Freq			χ^2
4.4	25	25	25	

(b)

Table 2: The intervals, class frequencies, and χ^2 values for attribute sepal-length after Phase 1 and Phase 2. The χ^2 thresholds are (a) 3.22 and (b) 50.6.

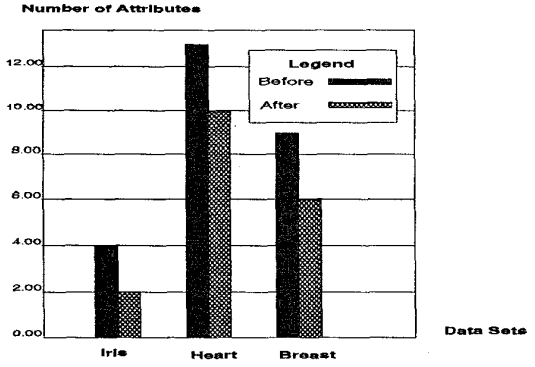

Figure 1: Number of attributes: original vs. those after Chi2 processing.

reduced from 4 to 2 (petal length and petal width), each has four values. For the breast cancer data, 3 attributes are removed from the original 9 attributes. The remaining 6 attributes have 3, 4, 4, 5, 3, and 3 discrete values respectively. For the heart disease data, the discrete attributes are left out in discretization and feature selection although they are used for consistency checking. Among the 5 continuous attributes (1, 4, 5, 8 and 10), only 2 attributes (5 and 8) should remain as suggested by Chi2, having 8 and 4 discrete values respectively. For the cancer and disease data sets, the default inconsistency rate is used, i.e., 0.

Second, we run C4.5 on both the original data sets and the dimensionally reduced ones. C4.5 is run using its default setting. Chi2 discretizes the training data and generates a mapping table, based on which the testing data are discretized.

Shown in Figure 2 are predictive accuracies and tree sizes of C4.5 for the three data sets. Predictive accuracy improves and tree size drops (by half) for the breast cancer and heart disease data. As for the Iris data, accuracy and tree size remain the same by using two attributes only (with 4 values each). In a way, it shows that C4.5 works pretty well without Chi2 for this data set.

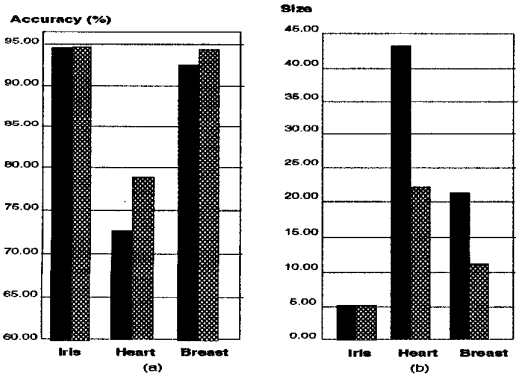

Figure 2: (a) Predictive accuracy and (b) size of decision trees of C4.5 for the three data sets after and before the Chi2 processing.

3.5 Empirical results on controlled data

The purpose of experimenting on the controlled data is to verify how effective Chi2 is in removing irrelevant attributes through discretizing numeric attributes. Therefore, it is only necessary to see if Chi2 can (1) discretize the relevant attribute(s) properly and (2) remove the irrelevant attributes.

Chi2 merged A_1 into three discrete values (1,2 and 3) corresponding to three classes (1,2, and 3); merged the other three attributes A_2, A_3, and A_4 into one value. That means that only A_1 should remain, and the noise (irrelevant) attributes should be removed.

For the modified Iris data, Chi2 merged six attributes out of eight. They are attributes 0, 1, 4, 5, 6 and 7. The first two are sepal-length and sepal-width. The last four are added noise (irrelevant) attributes. The remaining two attributes have been merged into 4 discrete values respectively as did in the real data experiment.

Through this set of controlled experiments, it is shown that Chi2 effectively discretizes numeric attributes and removes irrelevant attributes.

4 Discussions

ChiMerge requires a user to specify a proper significance level (α) which is used for merging values of all the attributes. No definite rule is given to choose this α. In other words, it is still a matter of trial-and-error, and clearly it is not easy to find a proper significance level for each problem. Phase 1 of Chi2 extends ChiMerge to an automated one. That is α is automatically varied until further merging is discontinued by the stopping criterion (the inconsistency rate). What makes Chi2 special is its capability of feature selection - a big step forward from discretization. In Phase 2 of Chi2, each attribute has its own significance level for merging in a round robin fashion. Merging stops when the inconsistency rate exceeds a specified one δ. This phase of Chi2 accomplishes feature selection. Another feature of Chi2 is that it can be applied to data with mixed attributes (e.g., Heart Disease Data). In addition, Chi2 can work with multi-calss data. This is an advantage over some statistic-based feature selection algorithms such as Relief [5] which is applicable only to the two-class data.

Other issues such as selecting δ, limitations of Chi2 as well as its computational complexity can be found in [6].

5 Conclusion

Chi2 is a simple and general algorithm that can automatically select a proper χ^2 value, determine the intervals of a numeric attribute, as well as select features according to the characteristics of the data. It guarantees that the fidelity of the training data can remain after Chi2 is applied. The empirical results on both the real data and controlled data have shown that Chi2 is a useful and reliable tool for discretization and feature selection of numeric attributes.

References

[1] H. Almuallim and T.G. Dietterich. Learning boolean concepts in the presence of many irrelevant features. *Artificial Intelligence*, 69(1-2):279–305, November 1994.

[2] J. Catlett. On changing continuous attributes into ordered discrete attributes. In *European Working Session on Learning*, 1991.

[3] U.M. Fayyad and K.B. Irani. The attribute selection problem in decision tree generation. In *AAAI-92, Proceedings Ninth National Conference on Artificial Intelligence*, pages 104–110. AAAI Press/The MIT Press, 1992.

[4] R. Kerber. Chimerge: Discretization of numeric attributes. In *AAAI-92, Proceedings Ninth National Conference on Artificial Intelligence*, pages 123–128. AAAI Press/The MIT Press, 1992.

[5] K. Kira and L.A. Rendell. The feature selection problem: Traditional methods and a new algorithm. In *AAAI-92, Proceedings Ninth National Conference on Artificial Intelligence*, pages 129–134. AAAI Press/The MIT Press, 1992.

[6] H. Liu and R. Setiono. Discretization of ordinal attributes and feature selection. Technical Report TRB4/95, Department of Info Sys and Comp Sci, National University of Singapore, April 1995, **"http: //www.iscs.nus.sg/liuh/chi2.ps"**.

[7] J.R. Quinlan. Induction of decision trees. *Machine Learning*, 1(1):81–106, 1986.

[8] J.R. Quinlan. *C4.5: Programs for Machine Learning*. Morgan Kaufmann, 1993.

[9] H. Ragavan and L. Rendell. Lookahead feature construction for learning hard concepts. In *Machine Learning: Proceedings of the Seventh International Conference*, pages 252–259. Morgan Kaufmann Pub. San Mateo, California, 1993.